

SaNDAl Global CEO Interview Preparation

Piter Garcia Bautista

Data and Business Analysis Track

September 16, 2026

Purpose. Prepare for a concise conversation with SaNDAl Global Founder and CEO Jeremy Williams Sr. The packet translates the submitted case assessment and verified experience into clear answers, evidence stories, and questions for the interview.

Main conclusion. Lead with practical judgment: understand the business and data, establish a baseline, choose the simplest reliable improvement, validate it, and scale only when the evidence supports it.

Private Zoom details remain in the calendar event and are intentionally excluded from this repository copy.

Interview Rapid Review

Three messages to land

1. **I understand the business and the data before choosing technology.**
2. **I improve one measurable, reversible part at a time.**
3. **I communicate what is known, what is inferred, and what still needs verification.**

Thirty second introduction

I am a data scientist, AI researcher, educator, and longtime systems engineer. I am strongest at turning an unclear business problem into a reliable workflow, analysis, or recommendation that technical and nontechnical people can use. I am interested in SaNDAl because the work combines real client responsibility with data, automation, AI, and community impact.

My core method

Understand the business → understand the data → establish a baseline → choose the simplest reliable improvement → build one bounded feature → test with users → measure → revise → scale only if justified

Important: Automation is broader than AI. A rule, form, script, API, integration, or dashboard may be the right answer.

Assessment in one minute

- Immediate risk: uncontrolled use of organizational information and unverified service answers.
- Best first improvement: acknowledge and route one website inquiry form.
- Why: measurable, useful, maintainable, and possible with ordinary automation.
- Privacy decision: executive authorization does not make every field or external tool appropriate.
- Crisis rule: stop harm first; contained, changed, and resolved are different states.

Best evidence stories

- **MLADIS:** source evidence → validated information → bounded use; human authorization remains separate; stale approval and uncertain sends are blocked.
- **Quantum and AI research:** Python validation, structured logging, experiment tracking, and reproducible comparison of large outputs.
- **Healthcare ML:** preprocessing, feature engineering, model experimentation, evaluation, and careful claim boundaries.
- **Teaching:** explain complex work in plain language, check understanding, and adapt without losing accuracy.

Strong phrases

- “The business goal should determine the technology.”
- “I would compare the improvement with an agreed baseline.”
- “I separate facts from assumptions before recommending action.”
- “I would offer a smaller permitted step that still moves the client toward the goal.”
- “A faster process is not better if it produces unreliable answers or moves work somewhere else.”
- “Contained is not the same as resolved.”

Questions for Jeremy

1. What distinguishes the people who become trusted with direct client responsibility at SaNDAI?
2. How are junior team members matched to projects, and who reviews work before it reaches a client?
3. What would strong performance look like on the first assignment?

Close

I am excited about the possibility of contributing. I bring deep technical and data experience, but I also understand that client work depends on judgment, communication, and follow-through. I would want my first project to make the client's work clearer, safer, and easier to act on.

Final reminders

- Join 10 minutes early using the private Zoom link in Calendar.
- Answer first; example second.
- Do not over-explain.
- Ask two questions.
- State interest clearly.

Full Interview Preparation Guide

Interview: SaNDAl Global interview with Founder and CEO Jeremy Williams Sr.

Calendar time: Wednesday, September 16, 2026, 5:00–5:30 PM Eastern

Format: Zoom; use the private link in the calendar event

Primary track: Data and Business Analysis

Secondary fit: AI subject matter work and project coordination

What this interview is likely testing

The written assessment already tested independent analysis, prioritization, risk recognition, executive communication, and judgment under incomplete information. This short CEO conversation is most likely checking whether I can explain that thinking clearly, represent SaNDAl professionally with clients, and contribute reliably in a project-based environment.

My goal is not to repeat every assessment answer. My goal is to make three things easy for Jeremy to see:

1. I understand business problems before choosing technology.
2. I can turn ambiguous work into a small, measurable, safe improvement.
3. I communicate limits honestly and still keep the work moving.

My opening answer

Tell me about yourself

I am completing graduate degrees in Data Science at RIT and Teaching Computer Science at the University of Rochester, and I bring more than ten years of engineering and data-systems experience. My strongest fit is data and business analysis: I enjoy understanding how a process works, organizing the data behind it, and turning the result into something people can actually use. My current research involves reliable evaluation workflows for quantum-network routing, and I also build and maintain operational systems such as MLADIS. Across those settings, I work the same way: understand the problem, establish a baseline, improve one bounded part, test it, and expand only when the evidence supports it. That is why SaNDAl's combination of real client work, data, automation, AI, and community impact interests me.

Shorter version if interrupted:

I am a data scientist, AI researcher, educator, and longtime systems engineer. I am strongest at turning an unclear business problem into a reliable workflow, analysis, or recommendation that technical and nontechnical people can use. I am especially interested in SaNDAl because the work combines real client responsibility with data, automation, AI, and community impact.

Why SaNDAl

I want work where the analysis has to survive contact with a real organization. SaNDAl's model gives emerging professionals direct client exposure and defined responsibility while working with experienced people. That fits me because I already bring substantial engineering and data experience, but I also want to keep growing in client discovery, executive communication, and responsible AI delivery. I also connect with the idea that technology should serve people and that data systems should tell the truth.

My operating approach in plain language

The core principle

Understand before automating, establish a baseline before changing, automate incrementally, validate every deployment, and scale only when the evidence supports it.

The process

1. **Understand the business and the people.** What outcome matters, how does the work happen now, and where is judgment required?
2. **Understand the data.** What exists, who owns it, how reliable is it, and what can appropriately be used?
3. **Establish a baseline.** How long does the current process take, where does it fail, and how will we know an improvement is real?
4. **Choose the simplest reliable improvement.** That might be a process change, rule, form, script, integration, dashboard, or AI.
5. **Build one bounded feature.** Keep it measurable and reversible.
6. **Test with the people who own the work.** Compare the result with the baseline and inspect errors and exceptions.
7. **Deploy carefully, measure, revise, and repeat.** Do not scale a change that has not been validated.

Automation is not the same as AI

I am positive about AI, but I do not begin with the assumption that AI is required. Automation can use workflow rules, scripts, APIs, forms, integrations, database triggers, or scheduled jobs. I use AI when it adds value beyond simpler methods—for example, when language variation, pattern recognition, or complex judgment support matters. The business goal should determine the technology, not the other way around.

Assessment debrief

What I noticed first

- The organization had outgrown parts of its operating structure, with limited staff capacity and decisions concentrated around the executive director.
- The digital journey had measurable friction: about one-third of registration attempts were incomplete, most traffic was mobile, and inquiries took about two business days to answer.

- Information was fragmented across forms, spreadsheets, email, and the donor system, without a common identifier.
- Staff were already using individual AI accounts without common rules for data handling, approved tools, output checks, or incident reporting.
- Funding was positive but needed resilience: a surplus existed, but grants represented 47% of revenue and only 17 active individual donors gave regularly.

My main recommendation

I would not start with a large AI transformation. I would put immediate data and service safeguards in place, then pilot acknowledgment and routing for one website inquiry form. We would map the current process, validate the baseline, assign an owner, implement the simplest workable solution, and compare response time, follow-up, errors, and staff effort. That gives the board a real result in three weeks without pretending the whole organization has been transformed.

The most important judgment call

Angela's authorization supported the business objective, but it did not automatically establish that every field in a mixed 4,500-record export was appropriate for donor targeting or safe to upload to an unspecified AI system. I would separate the donor information needed for the purpose, verify permitted use and system safeguards, check quality, and compare any model with a simpler giving-history approach. Permission to analyze also would not automatically authorize outreach.

Incident response

If a public assistant falsely promised rental assistance, I would first stop further harm. If the affected function could be isolated reliably, I would disable that function and provide verified information or a staff contact; if not, I would temporarily disable the assistant. Then I would preserve evidence, notify the responsible people, determine cause and scope, correct the problem, test the original question and variations, and restore service only after the responsible owners validated it. Contained, changed, and resolved are three different states.

Evidence stories to use

Story 1 MLADIS and safe operational automation

Situation: Operational information and customer interactions came from multiple systems and could not safely be treated as equally reliable.

Task: Build a workflow that could support people without allowing generated output or inferred signals to override source-backed facts.

Action: I organized the information lifecycle from raw evidence through cleaning, preparation, learning, and bounded use. I kept action behind a separate authorization boundary, required staff approval for customer responses, invalidated stale authorization when a new message or rule change appeared, and recorded independent verification states instead of calling the whole system complete.

Result: The system supports useful AI-assisted drafts while blocking escalation topics, duplicate or uncertain sends, and unsupported claims. It also makes failures traceable to a specific layer rather than hiding them under “the AI was wrong.”

What it proves: I can combine AI, ordinary automation, privacy, human oversight, testing, and rollback in one operational system.

Story 2 Reproducible quantum and AI research evaluation

Situation: Large experimental outputs from adaptive quantum-network routing needed to be compared reliably.

Task: Turn heterogeneous results into evidence that collaborators could review and reuse.

Action: I built Python workflows with shared datasets, automated validation, structured logging, experiment tracking, and comparison-ready reporting.

Result: The workflow makes assumptions, settings, and outcomes easier to reproduce and discuss rather than relying on isolated experiment outputs.

What it proves: I can structure ambiguous technical work, validate the data, and communicate results without overstating them.

Story 3 Healthcare machine learning

Situation: Healthcare data required careful preprocessing and evaluation before model results could be trusted.

Task: Develop machine-learning workflows that supported experimentation without confusing a model result with a clinical conclusion.

Action: I worked across preprocessing, feature engineering, model experimentation, evaluation, and AWS-oriented data preparation.

Result: The work supported evidence-based model comparison while keeping the limits of the analysis visible.

What it proves: I understand that high-stakes data work requires quality controls, context, and careful claims.

Story 4 Teaching and client communication

Situation: Learners and nontechnical collaborators do not benefit from jargon or a technically correct explanation they cannot use.

Task: Make complex ideas understandable while preserving accuracy.

Action: Through K–12 computer-science preparation and teaching placement, I break work into visible steps, use examples, check understanding, and adjust the explanation without changing the underlying concept.

Result: I can communicate technical tradeoffs to different audiences and create materials people can return to later.

What it proves: I can represent SaNDAl with clients, executives, staff, and community stakeholders.

Likely questions and answer maps

Why should we choose you

- More than ten years of engineering and data-systems experience.
- Current data-science and AI research, not only coursework.
- Strongest at requirements clarification, reliable workflows, validation, documentation, and translation between technical and business audiences.
- Practical judgment: I do not force AI where ordinary automation is better.
- Comfortable saying what is known, what is inferred, and what still needs verification.

How do you handle an ambiguous client request

I first restate the outcome, constraints, decision owner, and deadline. Then I walk through the current process with the people doing the work and inspect the data behind it. I separate facts from assumptions, establish a baseline, and propose the smallest useful improvement with a test and rollback plan. I keep the client informed about what is confirmed, what remains open, and what decision is needed.

What if a client insists on using AI immediately

I would not answer with a reflexive yes or no. I would clarify the business goal and identify the simplest safe path that produces value. If an AI approach is justified, I would constrain the data and use case, establish review and escalation, test against approved information, and measure it against a simpler baseline. If those conditions are not ready, I would offer a smaller permitted step that still moves the client toward the goal.

How do you prioritize when everything feels urgent

I prioritize immediate harm and trust first, then the smallest improvement that unlocks measurable progress, then the reporting needed for a decision. I also ask what must be true in three weeks rather than trying to solve the entire organization at once.

How do you know whether an automation worked

I compare it with an agreed baseline. The measure depends on the work: response time, documented follow-up, error rate, staff effort, completion rate, or another operational outcome. I also inspect exceptions and unintended effects. A faster process is not an improvement if it produces unreliable answers or creates more work elsewhere.

Tell me about a failure or disagreement

Use this structure rather than inventing a perfect story:

1. State the disagreement or failure plainly.
2. Explain what evidence was missing or what assumption broke.
3. Describe how I contained the impact and involved the right owner.
4. Explain the process or control added afterward.
5. End with what changed, not with blame.

How will you balance the program with school and teaching

I can meet the program's requirement of at least 20 hours per month for remote assignments. I plan work around fixed graduate classes and teaching responsibilities, agree on milestones and meeting requirements early, and protect focused project blocks so the work is distributed rather than left to one long session. I communicate conflicts early and do not accept a commitment I cannot meet.

What kind of work do you want

My strongest fit is Data and Business Analysis, especially discovery, requirements, data-quality review, workflow design, measurement, and executive-ready recommendations. I can also contribute to AI-focused work where the use case, data, evaluation, and human-review boundary are clear. I am open to project coordination when it supports delivery rather than replacing the analysis.

Questions to ask Jeremy

Choose two or three depending on time.

1. **What distinguishes the people who become trusted with direct client responsibility at SaNDAl?**
2. **How are junior team members matched to projects, and who reviews their work before it reaches a client?**
3. **What would strong performance look like in the first client assignment?**
4. **How does SaNDAl decide when a client problem needs AI versus a simpler automation or process change?**
5. **How do project hours, deadlines, meetings, and compensation work when assignments vary?**
6. **What kinds of nonprofit or community-impact projects are most active now?**

Closing answer

I would be excited to contribute. I bring deep technical and data experience, but I also understand that client work depends on judgment, communication, and follow-through. I am comfortable entering an unclear situation, learning how the organization actually works, and delivering one reliable improvement at a time. If selected, I would want my first project to make the client feel that the work became clearer, safer, and easier to act on.

Day of checklist

- Re-read the one-page quick sheet, not the entire packet.
- Open the calendar event and Zoom link 10 minutes early.
- Keep the assessment PDF and this guide available, but do not read answers word for word.
- Have water, headphones, charger, and a quiet background ready.
- Keep three phrases visible: **understand before automating, simplest reliable method, contained is not resolved.**
- Answer the question first, then give the example.
- Pause after the answer; do not fill every silence.
- Ask two thoughtful questions.
- End by expressing interest clearly.

Evidence and boundaries

- The interview title, date, time, format, and CEO role were verified from the live calendar and SaNDAl's official leadership page on September 16, 2026.
- The assessment discussion is derived from Piter's submitted client-readiness assessment PDF and the fictional Community Forward Alliance briefing.
- Experience claims are limited to the submitted SaNDAl resume and current MLADIS documentation.
- The Zoom URL remains private and is not copied into this public repository.
- Project availability is not guaranteed. The application committed to at least 20 hours per month for remote assignments, not a weekly maximum.

Source map

- SaNDAI application package README README.md
- SaNDAI application responses answers.md
- SaNDAI official founders page <https://sandaiglobal.net/meet-the-founders/>
- SaNDAI official junior team page <https://sandaiglobal.net/junior-team/>
- Private local source: SaNDAI Client-Readiness Proficiency Assessment.pdf
- Private Google Drive source: Student AI & Analytics Talent Program - Client Briefing File.pdf
- MLADIS local source: docs/data-store/pyramid-data-lifecycle.md
- MLADIS local source: docs/qa/feature-verification-ledger.md
- MLADIS local source: docs/qa/airbnb-response-workflow-verification.md